

MOCTAR SANKARA

DATA ENGINEER | DISTRIBUTED DATA PROCESSING | AI-ASSISTED DATA PLATFORMS

Washington, DC Metropolitan Area | +1 (415) 568-1997 | sankaramoctar@gmail.com | [GitHub](#) | [LinkedIn](#)

Willing to relocate to anywhere within the US (location agnostic)

PROFESSIONAL SUMMARY

Data Engineer with 7 years of professional experience across data engineering, data analytics, product analytics, and financial technology, including 4 years designing, building, and operating production data pipelines and analytical platforms. Experienced with Python, SQL, Spark, PySpark, Amazon EMR, Amazon S3, Airflow, dbt, Snowflake, and Jenkins. Owned distributed workflows processing approximately 100 TB of data and supported data-quality and SLA-monitoring environments containing approximately 100 million records. Experienced in end-to-end data platform delivery, dimensional modeling, AI-assisted development, automated validation, metadata management, production troubleshooting, performance optimization, technical documentation, and cross-functional partnership.

TECHNICAL SKILLS

Programming and Query Languages: Python, SQL, PySpark, Bash, LookML

Data Engineering: ETL/ELT, Batch Processing, Distributed Data Processing, Automated Data Pipelines, API Ingestion, Incremental Processing, Schema Evolution

Cloud & Data Platforms: Amazon EMR, Amazon S3, AWS IAM, AWS Secrets Manager, Snowflake, PostgreSQL

Data Modeling & Quality: Dimensional Modeling, Star Schemas, Fact and Dimension Tables, dbt Tests, Source-to-Target Validation, SLA Monitoring, Data Freshness, Root-Cause Analysis

Orchestration & DevOps: Airflow, dbt Core, Jenkins, Git, GitLab CI/CD, Code Reviews, Unit and Integration Testing, Poetry

AI & Analytics: Claude Code, LangChain, OpenAI, Streamlit, AI-Ready Metadata, NL-to-SQL, Looker, Tableau, Power BI, Superset, Metabase

PROFESSIONAL EXPERIENCE

BitSight Technologies Inc. | Boston, MA

Data Engineer, AI and Machine Learning Team | September 2025 – Present

- Owned a Jenkins-orchestrated ETL workflow using Bash, Amazon EMR, Spark, PySpark, and Amazon S3 to process approximately 100 TB of data and deliver curated products to internal and external customers.
- Developed an AI-assisted Airflow DAG generator with Claude Code, reducing initial DAG development from 3–4 hours to minutes while preserving testing, code review, and human approval controls.
- Built an AI-assisted dbt workflow for source analysis, model generation, testing, and documentation, reducing initial development from approximately one week to 1–2 hours with review and deployment controls.
- Built a governed AI-ready metadata platform covering approximately 12 tables and their relationships; implemented GitLab review and human approval, then distributed approved metadata as a reusable Python package.
- Developed a guarded NL-to-SQL application using OpenAI, LangChain, Streamlit, Snowflake, and AWS Secrets Manager; achieved 95–100% accuracy across an approximately 200-query evaluation while enforcing read-only access.
- Built incremental dbt models and analytical datasets for annual rating-methodology changes, enabling repeatable monthly analysis.
- Engineered a historical model that detected and reconciled added, removed, and renamed fields across methodology versions.
- Mentored a Data Science intern for more than three months across SQL, dbt, Airflow, S3-to-Snowflake architecture, validation, and Looker/LookML dashboard delivery.

Data Analyst | September 2022–August 2025

- Built and operated a daily Airflow pipeline from PostgreSQL to Snowflake, applying SQL and dbt transformations, dimensional models, tests, documentation, and LookML to tens of millions of records.
- Created data-quality and SLA-monitoring datasets spanning approximately 100 million records to track execution, duration, freshness, volume, completion, and attribution.

- Implemented schema validation, row-count reconciliation, source-to-target checks, dbt tests, logging, and reasonableness checks to detect missing, duplicated, and inconsistent data.
- Investigated pipeline failures and performance issues across PostgreSQL, Airflow, Jenkins, Snowflake, and dashboards; performed root-cause analysis, corrected logic, and reran affected workflows.
- Delivered a rating-preview pipeline from Amazon S3 through Snowflake and dbt to eight dashboards, supporting approximately 10,000 companies during a three-month test cycle.
- Built 10 Looker/LookML dashboards for operations, data quality, attribution, rating analysis, and stakeholder reporting.
- Developed reusable Python packages for AWS authentication, IAM role assumption, and entity resolution, supported by Poetry and GitLab CI/CD.

Pod Foods Co. | San Francisco, CA

Data Analyst, Product | November 2021 – September 2022

- Built Tableau, Power BI, and Apache Superset dashboards providing management with product, customer, financial, and operational performance insights.
- Partnered with product, engineering, operations, and business stakeholders to translate requirements into analytical datasets, KPIs, dashboards, and recommendations.

Financial Technology Analyst | October 2019 – October 2021

- Developed financial models and reporting for forecasting, budgeting, resource allocation, project tracking, and month-end analysis.
- Defined requirements and validation rules for invoicing, financial reporting, payment processing, and operational applications supporting more than 250 clients

EDUCATION

Hult International Business School | San Francisco, CA

Master of Science in International Business | 2019

University of Ghana Business School | Accra, Ghana

Master of Business Administration in Finance | 2018

Ecole de Commerce et de Gestion | Abidjan, Côte d'Ivoire

Bachelor of Science in Finance and Accounting | 2016

CERTIFICATION

Google Data Analytics Professional Certificate | Google | 2022